

## ▼ Title of the Project: Bank customer churn model

### Problem Statement:

1. Despite the continuous efforts of banks to attract and retain customers, the banking industry faces a persistent challenge in the form of customer churn, leading to financial losses and reduced customer satisfaction.

### Project Objective:

1. The aim of this project is to analyze the customer churn rate for bank because it is useful to understand why the customers leave.
2. After Analyzing we need to train a Machine Learning Model which can find the key factors that significantly influence the customer churn or attrition.
3. In the end will choose the most reliable model that will attach a probability to the churn to make it easier for customer service to target right customer in order to minimize their efforts to prevent customers churn.

1.Label Data Encoding 2.Features Scalling 3.Handling Imbalance Data a) Random Under sampling b) Random Over Sampling 4.Support vector machine classifier 5.Grid search for Hyperparameter Tunning

### Project Overview:

1. Churn refers to customers leaving a bank or discontinuing their banking services.
2. Banking Churn Analysis is a process of studying customer behavior in the banking industry to predict and understand customer attrition or churn.
3. Banking Churn Modeling aims to identify patterns and factors that contribute to customer churn, enabling banks to take proactive measures to retain customers and improve customer satisfaction.

About dataset: The bank customer churn dataset is a commonly used dataset for predicting customer churn in the banking industry. It contains information on bank customers who either left the bank or continue to be a customer.

## ▼ Importing Required Libraries

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

## ▼ Loading Dataset

```
df = pd.read_csv("https://github.com/YBI-Foundation/Dataset/raw/main/Bank%20Churn%20Modelling.csv")
```

## ▼ Describe/ Analysis Data

```
df.head()
```

	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	Num Produ
0	15634602	Hargrave	619	France	Female	42	2	0.00	
1	15647311	Hill	608	Spain	Female	41	1	83807.86	
2	15619304	Onio	502	France	Female	42	8	159660.80	
3	15704354	Boni	600	France	Female	30	1	0.00	

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 13 columns):
#   Column              Non-Null Count  Dtype
# 0  CustomerId          10000 non-null  int64
# 1  Surname              10000 non-null  object
# 2  CreditScore          10000 non-null  int64
# 3  Geography            10000 non-null  object
# 4  Gender               10000 non-null  object
# 5  Age                  10000 non-null  int64
# 6  Tenure               10000 non-null  int64
# 7  Balance              10000 non-null  float64
# 8  NumProdu             10000 non-null  int64
# 9  Churn                10000 non-null  bool
# 10  IsActive             10000 non-null  bool
# 11  NumProduct           10000 non-null  int64
# 12  NumProduct           10000 non-null  int64
```

```

---  -----
0   CustomerId      10000 non-null int64
1   Surname         10000 non-null object
2   CreditScore      10000 non-null int64
3   Geography        10000 non-null object
4   Gender           10000 non-null object
5   Age             10000 non-null int64
6   Tenure           10000 non-null int64
7   Balance          10000 non-null float64
8   Num Of Products  10000 non-null int64
9   Has Credit Card  10000 non-null int64
10  Is Active Member 10000 non-null int64
11  Estimated Salary 10000 non-null float64
12  Churn            10000 non-null int64
dtypes: float64(2), int64(8), object(3)
memory usage: 1015.8+ KB

```

```
df.duplicated('CustomerId').sum()
```

```
0
```

```
df= df.set_index('CustomerId')
```

```
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 10000 entries, 15634602 to 15628319
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   Surname                10000 non-null object
1   CreditScore            10000 non-null int64
2   Geography              10000 non-null object
3   Gender                 10000 non-null object
4   Age                   10000 non-null int64
5   Tenure                 10000 non-null int64
6   Balance                10000 non-null float64
7   Num Of Products        10000 non-null int64
8   Has Credit Card        10000 non-null int64
9   Is Active Member       10000 non-null int64
10  Estimated Salary       10000 non-null float64
11  Churn                  10000 non-null int64
dtypes: float64(2), int64(7), object(3)
memory usage: 1015.6+ KB

```

## ▼ Label Encoding

converting into catogorical values

```
df['Geography'].value_counts()
```

```

France    5014
Germany   2509
Spain     2477
Name: Geography, dtype: int64

```

```
df.replace({'Geography': {'France':2, 'Germany':1, 'Spain':0}},inplace=True)
```

```
df['Gender'].value_counts()
```

```

Male      5457
Female    4543
Name: Gender, dtype: int64

```

```
df.replace({'Gender': {'Male':0,'Female':1}},inplace=True)
```

```
df.replace({'Num Of Products': {1:0,2:1, 3:1, 4:1}},inplace=True)
```

#1-have the credit cards

#0-dont have the credit cards

```
df['Has Credit Card'].value_counts()
```

```

1      7055
0      2945
Name: Has Credit Card, dtype: int64

```

```
#1- active members
#0- not active members
df['Is Active Member'].value_counts()

1    5151
0    4849
Name: Is Active Member, dtype: int64
```


## ▼ Feature Engineering

```
#1- zero bank balance left the bank
#0- zero bank balance still in bank
df.loc[(df['Balance']==0), 'Churn'].value_counts()

0    3117
1     500
Name: Churn, dtype: int64
```

```
df['Zero Balance'] = np.where(df['Balance']>0,1,0)
```

```
df['Zero Balance'].hist()
```



```
df.groupby(['Churn', 'Geography']).count()
```

	Surname	CreditScore	Gender	Age	Tenure	Balance	Num Of Products	Has Credit Card	Is Active Member	Estimated Salary	Z Bala
Churn	Geography										
0	0	2064	2064	2064	2064	2064	2064	2064	2064	2064	2
	1	1695	1695	1695	1695	1695	1695	1695	1695	1695	1
	2	4204	4204	4204	4204	4204	4204	4204	4204	4204	4
1	0	413	413	413	413	413	413	413	413	413	
	1	814	814	814	814	814	814	814	814	814	

## ▼ Define Label and features

```
df.columns
```

```
Index(['Surname', 'CreditScore', 'Geography', 'Gender', 'Age', 'Tenure',
      'Balance', 'Num Of Products', 'Has Credit Card', 'Is Active Member',
      'Estimated Salary', 'Churn', 'Zero Balance'],
      dtype='object')
```

```
x= df.drop(['Surname','Churn'],axis=1)
```

```
y =df['Churn']

x.shape, y.shape

((10000, 11), (10000,))
```

## ▼ Handling Imbalance Data


### 1.Undersampling

### 2.Oversampling

```
#1-left
#0- not left
df['Churn'].value_counts()

0    7963
1    2037
Name: Churn, dtype: int64
```

```
sns.countplot(x='Churn', data=df);
```



```
x.shape, y.shape

((10000, 11), (10000,))
```

## ▼ Random under sampling

```
from imblearn.under_sampling import RandomUnderSampler
```

```
rus= RandomUnderSampler(random_state=2529)
```

```
x_rus, y_rus= rus.fit_resample(x,y)
```

```
x_rus.shape, y_rus.shape, x.shape, y.shape
```

```
((4074, 11), (4074,), (10000, 11), (10000,))
```

```
#1-left
#0- not left
y.value_counts()
```


```
0    7963
1    2037
Name: Churn, dtype: int64
```

```
#after random under sampling
#1-left
#0- not left
y_rus.value_counts()

0    2037
1    2037
Name: Churn, dtype: int64
```

```
y_rus.plot(kind= 'hist')
```

<Axes: ylabel='Frequency'>



## ▼ Random Over Sampling

```
from imblearn.over_sampling import RandomOverSampler
```

```
ros= RandomOverSampler(random_state=2529)
```

```
x_ros, y_ros= ros.fit_resample(x,y)
```

```
x_ros.shape, y_ros.shape, x.shape, y.shape
```

```
((15926, 11), (15926,), (10000, 11), (10000,))
```

```
#1-left
#0- not left
y.value_counts()
```

```
0    7963
1    2037
Name: Churn, dtype: int64
```

```
#after random over sampling
```

```
#1-left
```

```
#0- not left
```

```
y_ros.value_counts()
```

```
1    7963
0    7963
Name: Churn, dtype: int64
```

```
y_ros.plot(kind='hist')
```



## ▼ Train Test Split Method

```
from sklearn.model_selection import train_test_split
```

## ▼ Split original Data

```
x_train, x_test, y_train, y_test = train_test_split(x,y, test_size=0.3, random_state=2529)
```

## ▼ Split Random Under sample Data

```
x_train_rus, x_test_rus, y_train_rus, y_test_rus = train_test_split(x_rus,y_rus, test_size=0.3, random_state=2529)
```

## ▼ Split Random over sample Data

```
x_train_ros, x_test_ros, y_train_ros, y_test_ros = train_test_split(x_ros,y_ros, test_size=0.3, random_state=2529)
```

## ▼ Standardize Features

### Data preprocessing

```
from sklearn.preprocessing import StandardScaler
```

```
sc = StandardScaler()
```

## ▼ Standardize original data

```
x_train[['CreditScore', 'Age', 'Tenure','Balance','Estimated Salary']] = sc.fit_transform(x_train [['CreditScore', 'Age', 'Tenure','Balance','Estimated Salary']])
```

```
x_test[['CreditScore', 'Age', 'Tenure','Balance','Estimated Salary']] = sc.fit_transform(x_test [['CreditScore', 'Age', 'Tenure','Balance','Estimated Salary']])
```

## ▼ Standardize Random under sample data

```
x_train_rus[['CreditScore', 'Age', 'Tenure','Balance','Estimated Salary']] = sc.fit_transform(x_train_rus [['CreditScore', 'Age', 'Tenure','Balance','Estimated Salary']])
```

```
x_test_rus[['CreditScore', 'Age', 'Tenure','Balance','Estimated Salary']] = sc.fit_transform(x_test_rus [['CreditScore', 'Age', 'Tenure','Balance','Estimated Salary']])
```

## ▼ Standardize Random over sample data

```
x_train_ros[['CreditScore', 'Age', 'Tenure','Balance','Estimated Salary']] = sc.fit_transform(x_train_ros [['CreditScore', 'Age', 'Tenure','Balance','Estimated Salary']])
```

```
x_test_ros[['CreditScore', 'Age', 'Tenure','Balance','Estimated Salary']] = sc.fit_transform(x_test_ros [['CreditScore', 'Age', 'Tenure'
```

## ▼ Support vector Classifier

```
from sklearn.svm import SVC
svc = SVC()
```

```
svc.fit(x_train, y_train)
```

```
▼ SVC
SVC()
```

```
y_pred= svc.predict(x_test)
```

## ▼ Model Accuracy

```
from sklearn.metrics import confusion_matrix, classification_report
```

```
confusion_matrix(y_test, y_pred)
```

```
array([[2381,  33],
       [ 436, 150]])
```

```
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.85	0.99	0.91	2414
1	0.82	0.26	0.39	586
accuracy			0.84	3000
macro avg	0.83	0.62	0.65	3000
weighted avg	0.84	0.84	0.81	3000

support vector machine accuracy is 0.84= 84%

recall is = 0.26 = 26%

## ▼ Hyperparameter tuning

### cross-validation

```
from sklearn.model_selection import GridSearchCV
from sklearn.svm import SVC
```

```
param_grid = {
    'C': [0.1, 1, 10],
    'gamma': [1, 0.1, 0.01],
    'kernel': ['rbf'],
    'class_weight': ['balanced']
}
```

```
grid = GridSearchCV(SVC(), param_grid, refit=True, verbose=2, cv=2)
grid.fit(x_train, y_train)
```

```

Fitting 2 folds for each of 9 candidates, totalling 18 fits
[CV] END ..C=0.1, class_weight=balanced, gamma=1, kernel=rbf; total time= 1.6s
[CV] END ..C=0.1, class_weight=balanced, gamma=1, kernel=rbf; total time= 2.3s
[CV] END C=0.1, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 1.8s
[CV] END C=0.1, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 1.2s
[CV] END C=0.1, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 1.2s
[CV] END C=0.1, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 1.3s
[CV] END ...C=1, class_weight=balanced, gamma=1, kernel=rbf; total time= 1.4s
[CV] END ...C=1, class_weight=balanced, gamma=1, kernel=rbf; total time= 1.4s
[CV] END ..C=1, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 1.0s
[CV] END ..C=1, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 1.0s
[CV] END .C=1, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 1.1s
[CV] END .C=1, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 1.9s

```

```
print(grid.best_estimator_)
```

```
SVC(C=10, class_weight='balanced', gamma=1)
```

```
[CV] END C=10, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 1.0s
```

```
grid_predictions = grid.predict(x_test)
```

```
! ..... xvs
```

```
confusion_matrix(y_test, grid_predictions)
```

```
array([[2159, 255],
       [ 343, 243]])
```

```
print(classification_report(y_test, grid_predictions))
```

	precision	recall	f1-score	support
0	0.86	0.89	0.88	2414
1	0.49	0.41	0.45	586
accuracy			0.80	3000
macro avg	0.68	0.65	0.66	3000
weighted avg	0.79	0.80	0.79	3000

accuracy = 0.80= 80% recall =0.41= 41%

## ▼ Model with random under sampling

```
svc_rus= SVC()
```

```
svc_rus.fit(x_train_rus, y_train_rus)
```

```
▼ SVC
SVC()
```

```
y_pred_rus =svc_rus.predict(x_test_rus)
```

## ▼ Model accuracy

```
confusion_matrix(y_test_rus, y_pred_rus)
```

```
array([[470, 157],
       [174, 422]])
```

```
print(classification_report(y_test_rus, y_pred_rus))
```

	precision	recall	f1-score	support
0	0.73	0.75	0.74	627
1	0.73	0.71	0.72	596
accuracy			0.73	1223
macro avg	0.73	0.73	0.73	1223
weighted avg	0.73	0.73	0.73	1223

accuracy = 0.73= 73% recall =0.71= 71%



```


from sklearn.model_selection import GridSearchCV
from sklearn.svm import SVC

param_grid = {
    'C': [0.1, 1, 10],
    'gamma': [1, 0.1, 0.01],
    'kernel': ['rbf'],
    'class_weight': ['balanced']
}

grid_rus = GridSearchCV(SVC(), param_grid, refit=True, verbose=2, cv=2)
grid_rus.fit(x_train_rus, y_train_rus)

Fitting 2 folds for each of 9 candidates, totalling 18 fits
[CV] END ..C=0.1, class_weight=balanced, gamma=1, kernel=rbf; total time= 0.2s
[CV] END ..C=0.1, class_weight=balanced, gamma=1, kernel=rbf; total time= 0.2s
[CV] END C=0.1, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 0.2s
[CV] END C=0.1, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 0.2s
[CV] END C=0.1, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 0.2s
[CV] END C=0.1, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 0.2s
[CV] END ...C=1, class_weight=balanced, gamma=1, kernel=rbf; total time= 0.3s
[CV] END ...C=1, class_weight=balanced, gamma=1, kernel=rbf; total time= 0.2s
[CV] END ..C=1, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 0.2s
[CV] END ..C=1, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 0.2s
[CV] END .C=1, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 0.2s
[CV] END .C=1, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 0.2s
[CV] END ...C=10, class_weight=balanced, gamma=1, kernel=rbf; total time= 0.2s
[CV] END ...C=10, class_weight=balanced, gamma=1, kernel=rbf; total time= 0.3s
[CV] END .C=10, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 0.2s
[CV] END .C=10, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 0.2s
[CV] END C=10, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 0.2s
[CV] END C=10, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 0.2s

```



```

print(grid_rus.best_estimator_)

SVC(C=1, class_weight='balanced', gamma=0.1)

grid_predictions_rus = grid_rus.predict (x_test_rus)

confusion_matrix(y_test_rus, y_pred_rus)

array([[470, 157],
       [174, 422]])

print(classification_report(y_test_rus, y_pred_rus))

```

	precision	recall	f1-score	support
0	0.73	0.75	0.74	627
1	0.73	0.71	0.72	596
accuracy			0.73	1223
macro avg	0.73	0.73	0.73	1223
weighted avg	0.73	0.73	0.73	1223

## Model with random over sampling

```

svc_ros= SVC()

svc_ros.fit(x_train_ros, y_train_ros)

```

```

graph TD
    A[SVC] --> B[SVC()]

```

```

y_pred_ros =svc_ros.predict(x_test_ros)

```

## Model accuracy

```

confusion_matrix(y_test_ros, y_pred_ros)

array([[1823,  556],
       [ 626, 1773]])

print(classification_report(y_test_ros, y_pred_ros))

```

	precision	recall	f1-score	support
0	0.74	0.77	0.76	2379
1	0.76	0.74	0.75	2399
accuracy			0.75	4778
macro avg	0.75	0.75	0.75	4778
weighted avg	0.75	0.75	0.75	4778

accuracy = 0.75= 75% recall =0.74= 74%

```


from sklearn.model_selection import GridSearchCV
from sklearn.svm import SVC

param_grid = {
    'C': [0.1, 1, 10],
    'gamma': [1, 0.1, 0.01],
    'kernel': ['rbf'],
    'class_weight': ['balanced']
}

grid_ros = GridSearchCV(SVC(), param_grid, refit=True, verbose=2, cv=2)
grid_ros.fit(x_train_ros, y_train_ros)

Fitting 2 folds for each of 9 candidates, totalling 18 fits
[CV] END ..C=0.1, class_weight=balanced, gamma=1, kernel=rbf; total time= 5.0s
[CV] END ..C=0.1, class_weight=balanced, gamma=1, kernel=rbf; total time= 8.7s
[CV] END C=0.1, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 3.3s
[CV] END C=0.1, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 5.3s
[CV] END C=0.1, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 7.4s
[CV] END C=0.1, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 4.9s
[CV] END ...C=1, class_weight=balanced, gamma=1, kernel=rbf; total time= 6.6s
[CV] END ...C=1, class_weight=balanced, gamma=1, kernel=rbf; total time= 3.2s
[CV] END ..C=1, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 2.5s
[CV] END ..C=1, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 3.0s
[CV] END .C=1, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 3.6s
[CV] END .C=1, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 2.6s
[CV] END ...C=10, class_weight=balanced, gamma=1, kernel=rbf; total time= 3.0s
[CV] END ...C=10, class_weight=balanced, gamma=1, kernel=rbf; total time= 3.4s
[CV] END .C=10, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 3.8s
[CV] END .C=10, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 2.8s
[CV] END C=10, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 2.7s
[CV] END C=10, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 2.6s

```



```

print(grid_ros.best_estimator_)

SVC(C=10, class_weight='balanced', gamma=1)

grid_predictions_ros = grid_ros.predict (x_test_ros)

confusion_matrix(y_test_ros, grid_predictions_ros)

array([[2047,  332],
       [  68, 2331]])

print(classification_report(y_test_ros, grid_predictions_ros))

```

	precision	recall	f1-score	support
0	0.97	0.86	0.91	2379
1	0.88	0.97	0.92	2399
accuracy			0.92	4778
macro avg	0.92	0.92	0.92	4778
weighted avg	0.92	0.92	0.92	4778

accuracy = 0.92= 92% recall =0.97= 97%

## ▼ Results:

### Normal Data

```
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.85	0.99	0.91	2414
1	0.82	0.26	0.39	586
accuracy			0.84	3000
macro avg	0.83	0.62	0.65	3000
weighted avg	0.84	0.84	0.81	3000

## ▼ After Hyperparameter tuning

```
print(classification_report(y_test, grid_predictions))
```

	precision	recall	f1-score	support
0	0.86	0.89	0.88	2414
1	0.49	0.41	0.45	586
accuracy			0.80	3000
macro avg	0.68	0.65	0.66	3000
weighted avg	0.79	0.80	0.79	3000

## ▼ Random under sampling

```
print(classification_report(y_test_rus, y_pred_rus))
```

	precision	recall	f1-score	support
0	0.73	0.75	0.74	627
1	0.73	0.71	0.72	596
accuracy			0.73	1223
macro avg	0.73	0.73	0.73	1223
weighted avg	0.73	0.73	0.73	1223

## ▼ After Hyperparameter tuning On Random under sampler

```
print(classification_report(y_test_rus, y_pred_rus))
```

	precision	recall	f1-score	support
0	0.73	0.75	0.74	627
1	0.73	0.71	0.72	596
accuracy			0.73	1223
macro avg	0.73	0.73	0.73	1223
weighted avg	0.73	0.73	0.73	1223

## ▼ Random over sampling

```
print(classification_report(y_test_ros, y_pred_ros))
```

	precision	recall	f1-score	support
0	0.74	0.77	0.76	2379
1	0.76	0.74	0.75	2399

accuracy			0.75	4778
macro avg	0.75	0.75	0.75	4778
weighted avg	0.75	0.75	0.75	4778

## ▼ After Hyperparameter tuning On Random over sampler

```
print(classification_report(y_test_ros, grid_predictions_ros))
```

	precision	recall	f1-score	support
0	0.97	0.86	0.91	2379
1	0.88	0.97	0.92	2399
accuracy			0.92	4778
macro avg	0.92	0.92	0.92	4778
weighted avg	0.92	0.92	0.92	4778

## Churn Prediction:

In this classification problem, a lot of information about the consumers was provided, and the dataset was pretty clean with no missing values and no duplicate values. With only two classes in the objective feature (0: not churned, 1: churned), it was a binary classification challenge. The classes were imbalanced and the models were predicting all 0s in the target feature. We had employed the over-sampling technique to address this. we found repeated categorical values in some of the columns, We used LabelEncoder() to assign numerical values for categorical values. Later on, the Age, Balance and Vintage columns were scaled using standard scalar by projecting mean to zero and variance to one.

Some key lessons that the Bank can focus to bring down the churn rate:-

- 1. Greater percentage of females are likely to churn.
- 2. Customers who has not done any transaction in the past 3 months are more likely to churn.
- 3. Customers with credit card are more in number and are more likely to churn.
- Customers with poor credit ratings dominate the dataset and are also more likely to exit.
- 5. High-income customers are difficult to keep as they are more prone to churn.